



# SC3020 DBMS Project 1



A comprehensive analysis of the database system implementation & design



Storage System



B+ Tree Index



Performance  
Analysis

Created by: Chong Qi En, Goh Jun Keat,  
Cheah Jun An Douglas, Chin Linn Sheng, Nah Wei Jie

# Table Of Contents




---

1. Project Overview
2. Storage Components
3. B+ Tree Implementation & Optimisation
4. Tasks 1-3
5. Conclusion

# Project Introduction

## Project Overview

This project focuses on designing and implementing two core DBMS components:

-  Disk-supported storage system with 4KB page size
-  B+ tree index for efficient data retrieval
-  Uses NBA game data for evaluation

## Project Context

The project simulates real DBMS behavior while keeping complexity limited to:

 **Single binary file** for data storage

 **Separate binary file** for indexing

## Project Components



### Storage System

Fixed 4096 byte blocks,  
compact 22-byte records



### B+ Tree Index

Efficient range queries on  
FT\_PCT\_home field



### Data Loader

Parses NBA data from text file  
into compact records



### Query Engine

Range queries and efficient  
record deletion



### NBA Game Data

26,651 games, 9 attributes per game

# Record Storage

## 22-Byte Record Format

### 9 Attributes

4 bytes

**FG\_PCT\_Home**

- stored as float

4 bytes

**FT\_PCT\_Home**

- stored as float

4 bytes

**FG3\_PCT\_Home**

- stored as float

4 bytes

**Team ID**

- stored as 32-bit integer

2 bytes

**Date**

- 16-bit int no. of days from 2000-01-01

1 byte

**PTS\_Home**

- stored as 8-bit integer

1 byte

**AST\_Home**

- stored as 8-bit integer

1 byte

**REB\_Home**

- stored as 8-bit integer

1 byte

**Home Team Wins**

- stored as 8-bit integer

### Record Structure

- ✓ Fixed Size: **22 bytes** with no internal padding
- ✓ Each record is contained within 1 single block
- ✓ Fixed-length fields used for efficient storage
- ✓ No record header
- ✓ Each field is represented with the minimum number of bits
- ✓ Field sequence reordered for space optimisation
- ✓ Optimised for 4KB block size

### Compact Storage Techniques

- ✚ Used pragma pack (push,1)

# Block Storage

## 4KB Block Structure




### Block Structure

- ✓ Fixed size: **4096 bytes** (4KB)
- ✓ Sequential packing of fixed-length records
- ✓ No page header, block header or slot directory
- ✓ Records may begin at any byte offset within a block
- ✓ Records are accessed via block ID and offset
- ✓ Blocks are stored contiguously on the disk
- ✓ Mimics modern file system block size



### Block Information

Block size (byte)	4096
Record Size (byte)	22
Max Records per Block	186

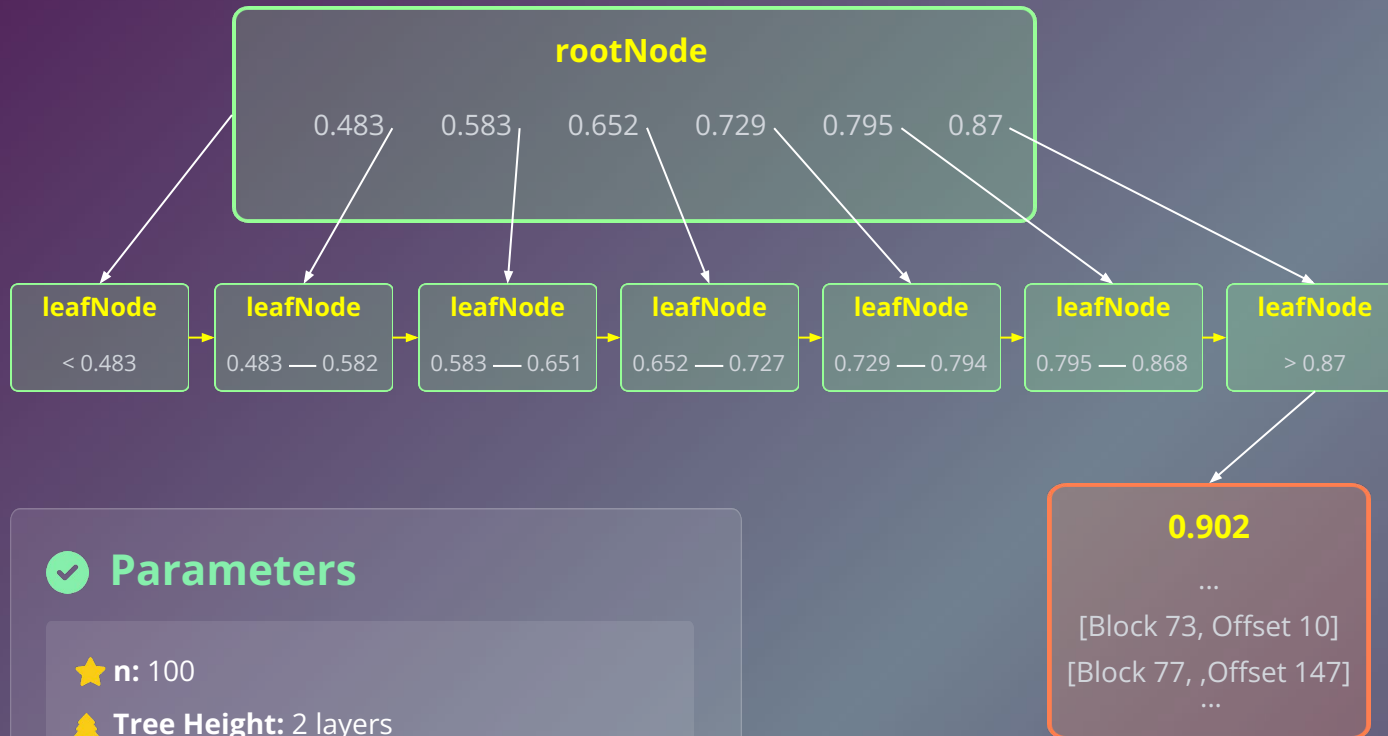


Block n (4 KB)

< 22 bytes

# B+ Tree Component Design

## B+ Tree Structure



### ✓ Parameters

- ★ **n:** 100
- 🌲 **Tree Height:** 2 layers
- 🔥 **Total nodes:** 8

### 🔑 RecordRef Structure

- ✓ Contains `block_id` and `record_offset`
- ✓ Uniquely identifies physical location of records
- ✓ Multiple RecordRef objects grouped under same key in leaf nodes
- ✓ Supports duplicate key entries without duplicating the key itself

### 🏠 BPlusNode Structure

- ✓ Unified structure for both leaf and internal nodes
- ✓ Internal nodes use key arrays as separators
- ✓ Internal nodes maintain pointers to guide search
- ✓ Leaf nodes store key-record associations as RecordRef objects
- ✓ Leaf nodes linked sequentially through `next_leaf` pointers

# B+ Tree Parameters & Optimization

## Parameter Calculations

### Block Size

4096 bytes

### Node Overhead

27 bytes

### Avg. Repeats

1.01466

### Conservative Repeats

2

## Optimal Parameters

★ **Optimal n:** 254

🌲 **Tree Height:** 2 layers

🍃 **Leaf nodes:** 104

\*Based on uniform distribution of keys

## n-Value Comparison

Parameter n	Tree Height	Leaf Nodes
n=50	3	526
n=100	3	263
n=200	2	132

### Key Considerations:

- ✓ Balance between tree height and node utilization
- ✓ n=100 provides good balance (3 layers, 263 leaf nodes)
- ✓ n=200 reduces tree height but increases branching factor

# B+ Tree Operations

## + Insert Operation

- ✓ Inserts key-recordRef pairs while maintaining tree balance
- ✓ Follows insertion algorithm from course slides
- ✓ Handles splitting of nodes when overflow occurs

## Insertion Algorithm

- 1 Search to find which leaf node to insert to
- 2 If leaf node is not full, insert it
- 3 Otherwise, split leaf node into 2 and distribute keys
- 4 Insert into parent and create new root if any
- 5 Repeat until a parent with no splits is required

## 🔍 Range Query Efficiency

B+ trees are optimized for range-based searches due to their linked leaf nodes structure.

```
</> searchGreaterThan(float key)
```

Returns all records with keys greater than the search key

## Range Query Algorithm

- 1 Start from root node
- 2 Follow appropriate pointer based on search key
- 3 Continue until reaching leaf node
- 4 Scan leaf node and follow pointer to next leaf block
- 5 Continue scanning leaf nodes until range limit



# Task 1 Results: Data Loading & Storage

## Data Storage Statistics



RECORD SIZE

**22 bytes**



TOTAL RECORDS

**26,651**



RECORDS PER BLOCK

**186**

max capacity



TOTAL BLOCKS

**144**

### Block Calculation Formula

```
ttlBlks = (ttlRecs + MAX_RECORDS_PER_BLOCK - 1) /  
          MAX_RECORDS_PER_BLOCK
```

### Block Structure Visualization



- Full blocks (186 recs)
- Last block (53 recs)

### ✓ Storage Efficiency

- ✚ Compact 22-byte record format minimizes space
- ✚ 186 records per 4KB block (optimal packing)
- ✚ Efficient use of space with minimal overhead

# Task 2 Results: B+ Tree Index Construction

## B+ Tree Statistics



Parameter n

**100**



Nodes

**8**



Levels

**2**



Construction Time

**13 ms**

## Root Node Keys

0.483

0.583

0.652

0.729

0.795

0.87

## Efficiency Highlights



- ✓ Compact 2-level structure with only 8 nodes
- ✓ Fast construction in just 13 milliseconds
- ✓ Optimal balance between depth and node capacity

# Task 3 Results: Record Deletion & Index Update

## Key Performance Metrics



Retrieval Time

**71.295 ms**

Running time of retrieval process



Records Deleted

**1778 games**

Records with FT\_PCT\_home > 0.9



Deletion Time

**9442 ms**

Total time for deleting 1,778 records



Index Node

**2 nodes**

Number of index nodes accessed

## Task 3 results visualization

```
=== Task 3 Results ===
Number of index nodes accessed: 2
Number of data blocks accessed: 144
Number of games deleted: 1778
Average FT_PCT_home of deleted records: 0.939637
Retrieval time (index + heap): 71.295 ms
Running time of deletion process: 9442 ms
```

## Deletion Summary



Average FT\_PCT\_home: 0.939637



B+ Tree Updates: 2 index nodes accessed

# Task 3 Results: Statistics of the updated B+ tree & Brute Force comparison

## Statistic of the B+ Tree Before and After

```
B+ tree loaded from disk: ft_pct_home.idx
=== B+ Tree Statistics ===
Parameter n: 100
Number of nodes: 8
Number of levels: 2
Root node keys: 0.483, 0.583, 0.652, 0.729, 0.795, 0.87
```

```
--- B+ Tree Statistics AFTER Deletion ---
=== B+ Tree Statistics ===
Parameter n: 100
Number of nodes: 7
Number of levels: 2
Root node keys: 0.483, 0.583, 0.652, 0.729, 0.795
B+ tree saved to disk: ft_pct_home.idx
```

### Explanation:

All deletions are at the last leaf ,  $>0.9$  is above 0.87

Since  $n = 100$  , the minimum number of keys must be between  $\lceil n/2 \rceil$  and  $n$  keys  $\rightarrow$  50 to 100 keys, After deletion of keys from  $>0.9$ , it reduces the number of key in the 0.87 root node to below 50, thus it performs a merge with 0.795 while still ensuring that the root node contains 50-100 keys.

## Performance Analysis

```
=== Brute-force Comparison ===
Linear scan would access: 144 data blocks
Linear scan estimated time: ~720 ms (estimated)
Assumption: 5 ms per block access and scan
Index speedup: ~10.0989x faster
```

### 💡 Key Insight

Matches are scattered across the heap, so both plans touch all 144 pages. The B+ tree still wins because it avoids a full record check path and coordinates access via the leaf range, yielding  $\sim 10\times$  shorter retrieval time under the fixed timing brute force access.

# Key Findings & Conclusions



## Efficient Space Utilization

Compact 22-byte records and 4KB blocks achieve near-ideal space utilization, storing 26,651 NBA game records in 144 blocks with minimal overhead.



## B+ Tree Performance

The B+ tree with  $n=100$  achieves 2-level structure with 8 nodes, providing efficient data retrieval.



## Record Deletion Efficiency


Successfully deleted 1,778 records with  $FT\_PCT\_home > 0.9$ , updating the B+ tree index while maintaining structural integrity.



## NBA Data Insights

The NBA dataset (2003-2022) shows a mean  $FT\_PCT\_home$  of 0.795 with a standard deviation of 0.114, with values ranging from 0.143 to 1.

B+ Tree vs. Linear Scan Performance



## Project Success

This project successfully implemented a disk-based storage system and B+ tree index, demonstrating efficient data storage and retrieval. The B+ tree structure reduced query time by over 10x compared to linear scanning, while maintaining structural integrity during record deletions.

**Thank You!**